MEET ZAVERI

DATA ENGINEER

+1(214) 836-7590 | meet.zaveri29@gmail.com | linkedin.com/in/meetzaveri | github.com/meetzaveri29

SUMMARY

Data Engineer building scalable data platforms across cloud and on-prem environments, with proven ability to reduce costs by 30%, cut execution time by 75%, and prevent $80K/month revenue leakage. Expert in designing distributed ETL/ELT pipelines on Azure and GCP using Python, PySpark, SQL, Airflow, and Databricks to power real-time analytics and ML infrastructure

Core Strengths: ETL/ELT Architecture • Cloud Data Warehousing • Data Optimization • Workflow Orchestration • Data Quality Frameworks • CI/CD Automation • Big Data Performance Engineering

TECHNICAL SKILLS

Languages & Frameworks: SQL , Python (Pandas, NumPy, PySpark), Shell Scripting, FastAPI
Big Data & Processing: Snowflake, dbt, Databricks , Apache Spark, Apache Airflow, Kafka, Hadoop (HDFS, Hive, Impala)
Cloud Platforms: Azure (Data Factory, Synapse, ADLS Gen2), GCP (BigQuery, Dataflow, Pub/Sub)
Databases: MS SQL Server , Teradata, PostgreSQL, MySQL, MongoDB
Data Engineering: ETL/ELT Design, Data Modeling (Star, Snowflake), Data Lakehouse, Data Quality, Performance Tuning
DevOps & Tools: Git, Github Actions, Jenkins, Bitbucket, Docker, Tableau, Power BI, JIRA, Confluence, Agile/Scrum

PROFESSIONAL EXPERIENCE

Data Engineer | Bank OZK | Dallas, TX — **January 2026 - Present**

- Developed 20+ dbt models on Snowflake for lending data mart migration from SQL Server, transforming raw lending data into analytics-ready layers and reducing average query execution time from 10 minutes to 30 seconds
- Streamlined release cycles by implementing GitHub Actions CI/CD pipelines for dbt and Snowflake deployments, automating environment-based promotions and pull request validations, significantly reducing manual deployment efforts

Data Engineer | ITMiracle | Dallas, TX — **September 2025 - January 2026**

- Architected scalable data pipelines using Azure Databricks and PySpark to process over 5 million records daily from REST APIs, streamlining data ingestion time by 65% through optimization and partitioning strategies
- Built FastAPI-based microservices for data delivery, serving real-time datasets to BI dashboards and ML feature stores with less than 100ms latency, improving downstream analytics accessibility by 50%
- Automated data quality frameworks with Python and SQL validation rules, achieving 99.5% data accuracy across multi-source integrations

Data Engineer | Accenture | Mumbai, India — **January 2021 - July 2023**

- Engineered distributed ETL pipelines for 3+ high-throughput revenue systems using Airflow, PySpark, and Informatica IICS, scaling ingestion for 10M+ records daily while improving reporting latency and Postal KPIs by 20%
- Developed end-to-end CI/CD automation with Git, Jenkins, and Bitbucket, eliminating manual deployments and preventing $80K/month revenue leakage caused by version control issues in production systems
- Designed and deployed real-time data ingestion POC of Teradata BTEQ using Pub/Sub → Dataflow → BigQuery, achieving 97% accuracy and informing enterprise migration from Teradata to GCP
- Optimized large-scale PySpark and Teradata workloads (1.5TB+ daily) by refactoring joins, implementing broadcast strategies, and tuning partition schemes - minimizing compute costs by 30% and execution time by 75%
- Automated batch monitoring and anomaly detection for 100+ SQL workflows (500M+ daily records) using Python + Splunk APIs, improving incident detection speed by 85% and achieving 99.9% SLA adherence
- Revamped data quality validation framework using Python, SQL, and Shell scripts across 20+ production pipelines, diminishing ingestion errors by 98% and standardizing exception handling
- Enforced GDPR-aligned archival and retention policies, improving data lifecycle reliability and reducing downtime by 15%

PROJECTS

Modern Data Platform (Azure & Databricks) - Built metadata-driven ELT pipeline across 10+ REST and relational sources using ADF, Databricks, and PySpark; implemented Bronze-Silver-Gold medallion architecture on ADLS with Synapse Dedicated SQL Pools, trimming Power BI refresh time from 2 hours to 5 minutes (96% improvement)

Fleet Analytics System (Hadoop & SAS) - Designed Spark + Hadoop pipeline processing 5GB+ daily telemetry logs; applied SAS regression modeling to predict mileage-related incidents, reducing violations by 15%

EDUCATION

The University of Texas at Dallas, Richardson, TX, USA — **August 2023 - May 2025**
Master of Science, Information Technology and Management, Business Intelligence and Analytics Track — **GPA 3.95/4.0**